

CS/ECE/ME532 Assignment 10

1. Neural net functions

- a) Sketch the function generated by the following 3-neuron ReLU neural network.


$$f(x) = 2(x - 0.5)_+ - 2(2x - 1)_+ + 4(0.5x - 2)_+$$

where $x \in \mathbb{R}$ and where $(z)_+ = \max(0, z)$ for any $z \in \mathbb{R}$. Note that this is a single-input, single-output function. Plot $f(x)$ vs x by hand.

- b) Consider the continuous function depicted below. Approximate this function with ReLU neural network with 2 neurons. The function should be in the form

$$f(x) = \sum_{j=1}^2 v_j (w_j x + b_j)_+$$

Indicate the weights and biases of each neuron and sketch the neural network function.



- c) A neural network f_w can be used for binary classification by predicting the label as $\hat{y} = \text{sign}(f_w(\mathbf{x}))$. Consider a setting where $\mathbf{x} \in \mathbb{R}^2$ and the desired classifier is -1 if both elements of \mathbf{x} are less than or equal to zero and $+1$ otherwise. Sketch the desired classification regions in the two-dimensional plane, and provide a formula for a ReLU network with 2-neurons that can produce the desired classification. For simplicity, assume in this questions that $\text{sign}(0) = -1$.
2. **Gradients of a neural net.** Consider a 2 layer neural network of the form $f(\mathbf{x}) = \sum_{j=1}^J v_j (\mathbf{w}_j^T \mathbf{x})_+$. Suppose we want to train our network on a dataset of N samples \mathbf{x}_i with corresponding labels y_i , using a least squares loss function $\mathcal{L} = \sum_{i=1}^n (f(\mathbf{x}_i) - y_i)^2$. Derive the gradient descent update steps for the input weights \mathbf{w}_j and output weights v_j .

- 3. Compressing neural nets.** Large neural network models can be approximated by considering low rank approximations to weight matrices. The neural network $f(\mathbf{x}) = \sum_{j=1}^J \mathbf{v}_j(\mathbf{w}_j^T \mathbf{x})_+$ can be written as

$$f(\mathbf{x}) = \mathbf{v}^T(\mathbf{W}\mathbf{x})_+.$$

where \mathbf{v} is a $J \times 1$ vector of the output weights and \mathbf{W} is a $J \times d$ matrix with i th row \mathbf{w}_j^T . Let $\sigma_1, \sigma_2, \dots$ denote the singular values of \mathbf{W} and assume that $\sigma_i \leq \epsilon$ for $i > r$. Let f_r denote the neural network obtained by replacing \mathbf{W} with its best rank r approximation $\hat{\mathbf{W}}_r$. Assuming that \mathbf{x} has unit norm, find an upper bound to the difference $\max_{\mathbf{x}} |f(\mathbf{x}) - f_r(\mathbf{x})|$. (Hint: for any pair of vectors \mathbf{a} and \mathbf{b} , the following inequality holds $\|\mathbf{a}_+ - \mathbf{b}_+\|_2 \leq \|\mathbf{a} - \mathbf{b}\|_2$).

- 4. Face Emotion Classification with a three layer neural network.** In this problem we return to the face emotion data studied previously. You may find it very helpful to use code from an activity (or libraries such as Keras and Tensorflow).

- a) Build a classifier using a full connected three layer neural network with logistic activation functions. Your network should
- take a vector $\mathbf{x} \in \mathbb{R}^{10}$ as input (nine features plus a constant offset),
 - have a single, fully connected hidden layer with 32 neurons
 - output a scalar \hat{y} .

Note that since the logistic activation function is always positive, your decision should be as follows: $\hat{y} > 0.5$ corresponds to a ‘happy’ face, while $\hat{y} \leq 0.5$ is not happy.

- b) Train your classifier using stochastic gradient descent (start with a step size of $\alpha = 0.05$) and create a plot with the number of epochs on the horizontal axis, and training accuracy on the vertical axis. Does your classifier achieve 0% training error? If so, how many epoch does it take for your classifier to achieve perfect classification on the training set?
- c) Find a more realistic estimate of the accuracy of your classifier by using 8-fold cross validation. Can you achieve perfect test accuracy?